

## Topic 2: Google Infrastructure and GFS

---

- What are the design principles behind the Google Infrastructure?
- Go into depth with GFS
- What kind of infrastructure does GFS target? What are the key assumptions and design goals behind it?
- Explain the architecture, consistency model, replication, fault tolerance
- What are its advantages and disadvantages?

Material: slides on Google infrastructure, GFS paper, "Web-search for a planet" paper

### Design Principles

Energy efficiency and price-performance ratio. (Web-search for a planet, page 1)

### GFS

GFS targets an infrastructure composed of many commodity hardware components.

GFS was created with the following assumptions:

- Failures are the norm
- Files are typically very large (multi-gigabyte)
- Large streaming reads (sequential) and small random reads (often buffered, also sequential)
- Write operations are typically append, they are typically large and sequential
- Support concurrent appending to a file with well-defined semantics
- High sustained bandwidth is more important than latency

### Architecture

The high-level architecture is composed of three actors: One master, many chunkservers, many clients. All run Linux.

The master acts as a DNS server for clients, i.e. it takes a request and returns an address to a chunkserver. Chunks are 64MB each, and have unique 64-bit identifiers: *chunk handles*.

### Consistency Model

The paper describes the model as a "relaxed consistency model". Garbage is collected lazily. It guarantees:

- Atomic file namespace mutations, such as file creation.
- A file region is:
  - *Consistent* if all clients see the same data from all replicas.
  - *Defined* if consistent and clients see their writes in their entirety.
- Records are appended atomically at least once somewhere
- Applications can handle inconsistent regions

## Replication

Replicas are spread across different racks. This makes them very resistant to faults, but increases the time it takes to write. The user defines a replication goal, which is typically 3.

Replicas are created for three reasons:

1. **Creation:** Initially, the replicas are empty
2. **Re-replication:** When replicas are lost, or replication goal redefined.
3. **Rebalancing:** Periodical rebalancing

Considerations: Resource utilization, recent creates, and separate racks.

## Fault Tolerance

1. Constant monitoring
2. Replicating crucial data
3. Fast and automatic recovery

## Conclusion

Only appropriate for very specific stacks. Does not support small files very well. High latency, especially with writes.